

R for ML

Applied Machine Learning with R

www.therbootcamp.com

@therbootcamp

January 2019

R is a programming language

From [Wikipedia](#) (emphasis added):

A programming language is a **formal language** that specifies a set of instructions that can be used to produce various kinds of output. Programming languages generally consist of **instructions for a computer**. Programming languages can be used to create programs that **implement specific algorithms**.

Algorithm

1. Load data
2. Extract variables
3. Run analysis
4. Print result

Implementation in R

```
#data <- read.table(link)
#variables <- data[,c('group', 'variable')]
#analysis <- lm(variable ~ group, data = variables)
#summary(analysis)
```

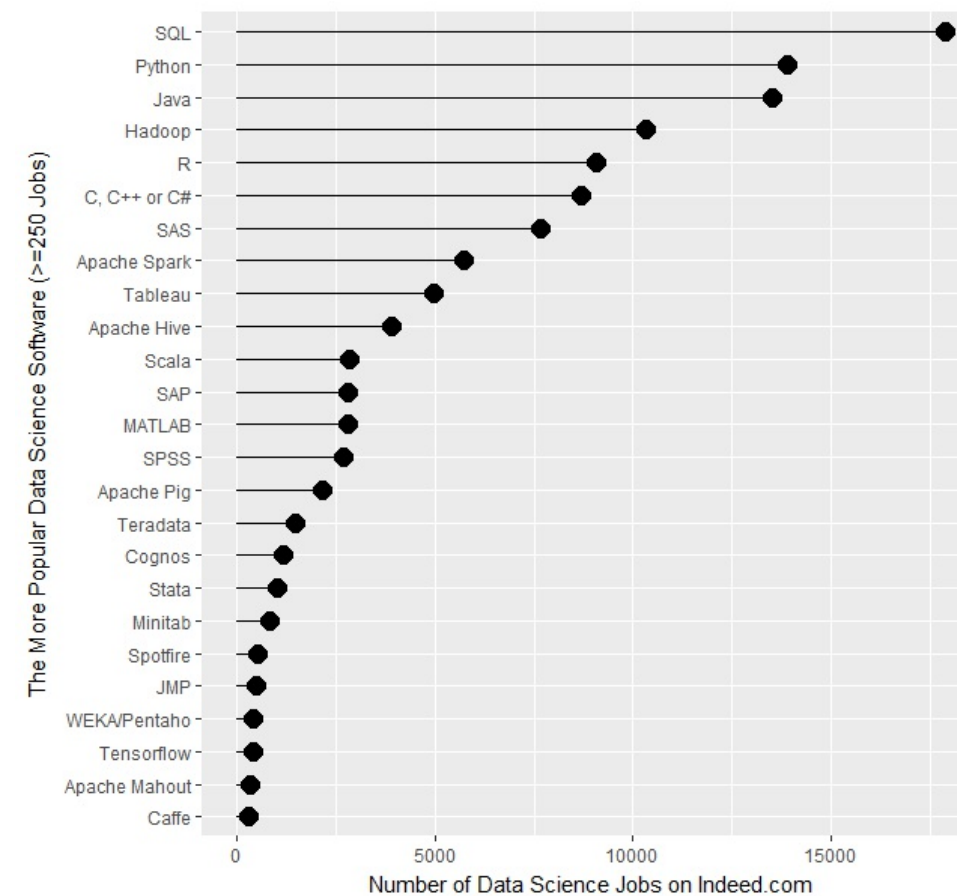
Why R?

R steadily **grows in popularity**.

Today, R is one of the **most popular languages for data science** and overall.

In terms of the number of data science jobs, **R beats SAS and Matlab**, and is on par with Python.

Image source: <https://i0.wp.com/r4stats.com/>



R is so popular because

There are many good reasons to prefer R over superficially more user friendly software such as **Excel** or **SPSS** or more complex programming languages like **C++** or **Python**.

Pro

1. **It's free**
2. Relatively **easy**
3. **Extensibility** (**CRAN**, packages)
4. **User base** (e.g., **stackoverflow**)
5. **Tidyverse** (dplyr, ggplot, etc.)
6. **RStudio**
7. **Productivity** options: **Latex**, **Markdown**, **GitHub**

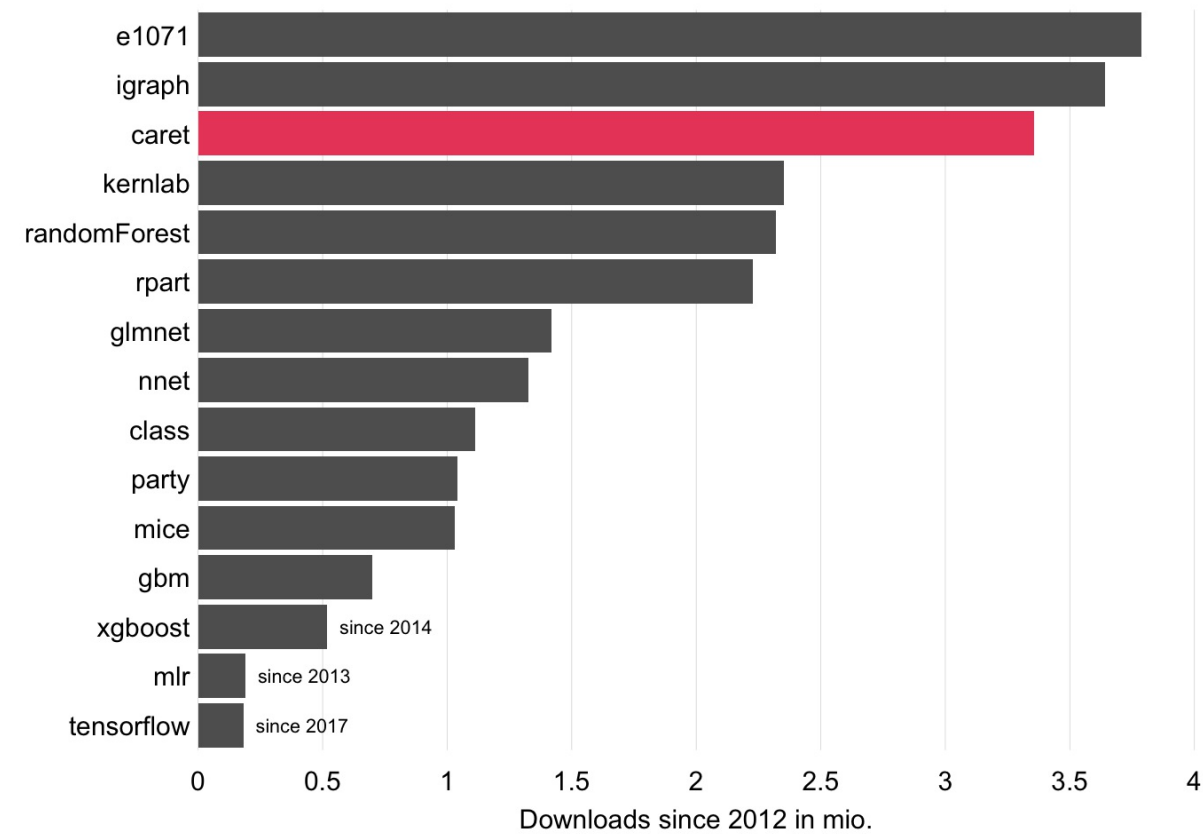
Con

It's slow, but...

Tidyverse Rcpp, **BH**: Links R to C++ and high-performance C++ libraries
rPython: Links R to Python
RHadoop: Links R to Hadoop for big data applications.

R is great for ML

...because of high-performance R packages (extensions) downloaded and used millions of times.



caret

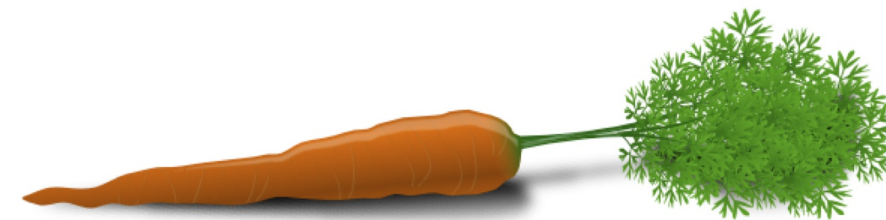
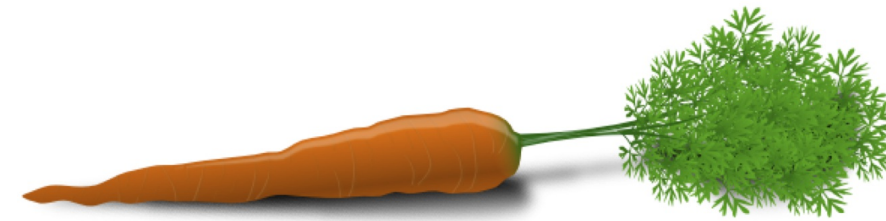
The **C**lassification **A**nd **R**egression **T**raining package is a meta-package to streamline the application of R's best machine learning tools.

caret facilitates...

- 1) **data pre-processing**
- 2) **feature selection**
- 3) **fitting, tuning, & model prediction**

Includes dozens of algorithms/models including...

regression, decision trees, random forests, neural nets, AdaBoost, elastic nets, lasso & ridge regression, support vector machines, etc.



10 basic R lessons

Essentials: 10 basic R lessons

1. Everything is an object
2. `<-` creates/changes objects
3. Everything happens through functions
4. Functions have (default) arguments
5. Find help with `?`
6. Data lives in data frames
7. 3 data types + factors
8. `formula` and `data` specify a model
9. Use RStudio and projects
10. Use editor, shortcuts, auto-complete

Essentials: 10 basic R lessons

1. **Everything is an object**
2. **<- creates/changes objects**
3. Everything happens through functions
4. Functions have (default) arguments
5. Find help with ?
6. Data lives in data frames
7. 3 data types + factors
8. formula and data specify a model
9. Use RStudio and projects
10. Use editor, shortcuts, auto-complete

```
# an object called some_name  
some_name <- c(1, 2, 3)  
  
# add 2 to the object's numbers  
some_name + 2
```

```
## [1] 3 4 5
```

```
# print object  
some_name
```

```
## [1] 1 2 3
```

```
# make change permanent  
some_name <- some_name + 2  
  
# print object  
some_name
```

```
## [1] 3 4 5
```

Essentials: 10 basic R lessons

1. Everything is an object
2. <- creates/changes objects
3. **Everything happens through functions**
4. Functions have (default) arguments
5. Find help with ?
6. Data lives in data frames
7. 3 data types + factors
8. formula and data specify a model
9. Use RStudio and projects
10. Use editor, shortcuts, auto-complete

```
# function c()
some_name <- c(1, 2, 3)

# function `+`()
some_name + 2
```

```
## [1] 3 4 5
```

```
# function print()
some_name
```

```
## [1] 1 2 3
```

```
# function mean()
mean(some_name)
```

```
## [1] 2
```

Essentials: 10 basic R lessons

1. Everything is an object
2. <- creates/changes objects
3. Everything happens through functions
4. **Functions have (default) arguments**
5. Find help with ?
6. Data lives in data frames
7. 3 data types + factors
8. formula and data specify a model
9. Use RStudio and projects
10. Use editor, shortcuts, auto-complete

```
# no argument  
mean()
```

```
## Error in mean.default(): argument "x" is missing, with no
```

```
# one (required) argument  
mean(x = c(1, 2, 3))
```

```
## [1] 2
```

```
# assume a missing value (NA)  
mean(x = c(1, 2, 3, NA))
```

```
## [1] NA
```

```
# changing default to handle NA  
mean(x = c(1, 2, 3, NA), na.rm = TRUE)
```

```
## [1] 2
```

Essentials: 10 basic R lessons

1. Everything is an object
2. <- creates/changes objects
3. Everything happens through functions
4. **Functions have (default) arguments**
5. Find help with ?
6. Data lives in data frames
7. 3 data types + factors
8. formula and data specify a model
9. Use RStudio and projects
10. Use editor, shortcuts, auto-complete

```
# mean with pipe %>%  
c(1, 2, 3) %>% mean()
```

```
## [1] 2
```

```
# mean with pipe %>% and NA  
c(1, 2, 3, NA) %>% mean()
```

```
## [1] NA
```

```
# changing default to handle NA  
c(1, 2, 3, NA) %>% mean(na.rm = TRUE)
```

```
## [1] 2
```

Essentials: 10 basic R lessons

1. Everything is an object
2. <- creates/changes objects
3. Everything happens through functions
4. Functions have (default) arguments
5. **Find help with ?**
6. Data lives in data frames
7. 3 data types + factors
8. formula and data specify a model
9. Use RStudio and projects
10. Use editor, shortcuts, auto-complete

?mean

mean {base}

R Documentation

Arithmetic Mean

Description

Generic function for the (trimmed) arithmetic mean.

Usage

```
mean(x, ...)
```

```
## Default S3 method:
```

```
mean(x, trim = 0, na.rm = FALSE, ...)
```

Arguments

x An R object. Currently there are methods for numeric/logical vectors and [date](#), [date-time](#) and [time interval](#) objects. Complex vectors are allowed for `trim = 0`, only.

trim the fraction (0 to 0.5) of observations to be trimmed from each end of `x` before the mean is computed. Values of `trim` outside that range are taken as the nearest endpoint.

na.rm a logical value indicating whether NA values should be stripped before the computation proceeds.

... further arguments passed to or from other methods.

Value

If `trim` is zero (the default), the arithmetic mean of the values in `x` is computed, as a numeric or complex vector of length one. If `x` is not logical (coerced to numeric), numeric (including integer) or complex, `NA_real_` is returned, with a warning.

If `trim` is non-zero, a symmetrically trimmed mean is computed with a fraction of `trim` observations deleted from each end before the mean is computed.

References

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

Essentials: 10 basic R lessons

1. Everything is an object
2. `<-` creates/changes objects
3. Everything happens through functions
4. Functions have (default) arguments
5. **Find help with ?**
6. Data lives in data frames
7. 3 data types + factors
8. `formula` and `data` specify a model
9. Use RStudio and projects
10. Use editor, shortcuts, auto-complete

?cor

cor {stats}

R Documentation

Correlation, Variance and Covariance (Matrices)

Description

`var`, `cov` and `cor` compute the variance of `x` and the covariance or correlation of `x` and `y` if these are vectors. If `x` and `y` are matrices then the covariances (or correlations) between the columns of `x` and the columns of `y` are computed.

`cov2cor` scales a covariance matrix into the corresponding correlation matrix *efficiently*.

Usage

```
var(x, y = NULL, na.rm = FALSE, use)
```

```
cov(x, y = NULL, use = "everything",  
    method = c("pearson", "kendall", "spearman"))
```

```
cor(x, y = NULL, use = "everything",  
    method = c("pearson", "kendall", "spearman"))
```

```
cov2cor(V)
```

Arguments

`x` a numeric vector, matrix or data frame.

`y` `NULL` (default) or a vector, matrix or data frame with compatible dimensions to `x`. The default is equivalent to `y = x` (but more efficient).

`na.rm` logical. Should missing values be removed?

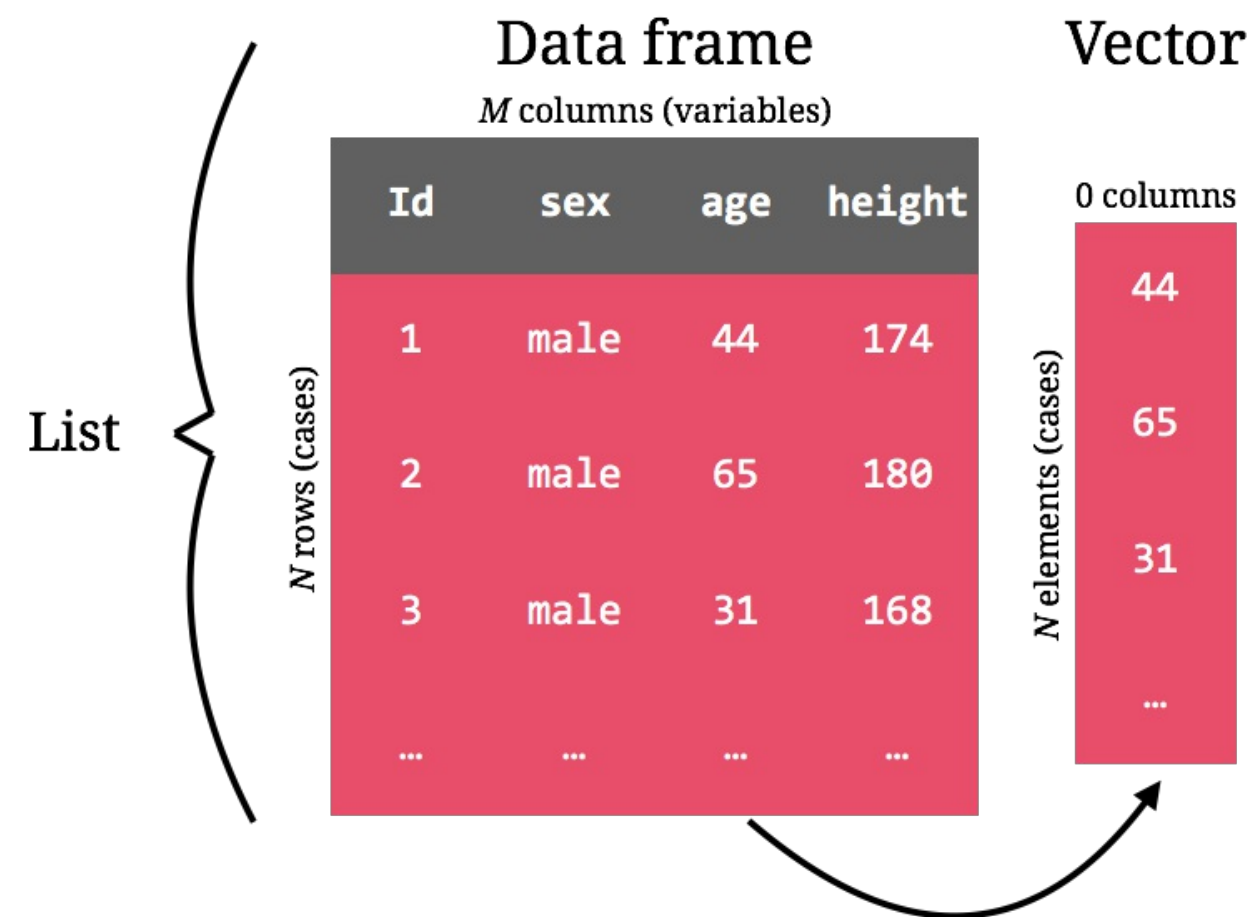
`use` an optional character string giving a method for computing covariances in the presence of missing values. This must be (an abbreviation of) one of the strings "everything", "all.obs", "complete.obs", "na.or.complete", or "pairwise.complete.obs".

`method` a character string indicating which correlation coefficient (or covariance) is to be computed. One of "pearson" (default), "kendall", or "spearman": can be abbreviated.

`V` symmetric numeric matrix, usually positive definite such as a covariance matrix.

Essentials: 10 basic R lessons

1. Everything is an object
2. `<-` creates/changes objects
3. Everything happens through functions
4. Functions have (default) arguments
5. Find help with `?`
6. **Data lives in data frames**
7. 3 data types + factors
8. `formula` and `data` specify a model
9. Use RStudio and projects
10. Use editor, shortcuts, auto-complete



Essentials: 10 basic R lessons

1. Everything is an object
2. `<-` creates/changes objects
3. Everything happens through functions
4. Functions have (default) arguments
5. Find help with `?`
6. Data lives in data frames
7. **3 data types + factors**
8. `formula` and `data` specify a model
9. Use RStudio and projects
10. Use editor, shortcuts, auto-complete

numeric Vector	character Vector	logical Vector
<code>.\$age</code>	<code>.\$sex</code>	<code>.\$sex=="male"</code>
44	"male"	TRUE
65	"female"	FALSE
31	"male"	TRUE
...

Essentials: 10 basic R lessons

1. Everything is an object
2. <- creates/changes objects
3. Everything happens through functions
4. Functions have (default) arguments
5. Find help with ?
6. Data lives in data frames
7. **3 data types + factors**
8. formula and data specify a model
9. Use RStudio and projects
10. Use editor, shortcuts, auto-complete

```
print(baselers)
```

```
## # A tibble: 10,000 x 20
##       id sex    age height weight income
##   <int> <chr> <int>   <dbl>   <dbl>   <dbl>
## 1     1 1 male    44    174.    113.    6300
## 2     2 2 male    65    180.    75.2   10900
## 3     3 3 fema... 31    168.    55.5    5100
## 4     4 4 male    27    209     93.8    4200
## 5     5 5 male    24    177.     NA     4000
## 6     6 6 male    63    187.    67.4   11400
## 7     7 7 male    71    152.    83.3   12000
## 8     8 8 fema... 41    156.    67.8    7600
## 9     9 9 male    43    176.    69.3    8500
## 10    10 10 fema... 31    166.    66.3    6100
## # ... with 9,990 more rows, and 14 more
## # variables
```

Essentials: 10 basic R lessons

1. Everything is an object
2. <- creates/changes objects
3. Everything happens through functions
4. Functions have (default) arguments
5. Find help with ?
6. Data lives in data frames
7. **3 data types + factors**
8. formula and data specify a model
9. Use RStudio and projects
10. Use editor, shortcuts, auto-complete

```
# select sex variable using $  
baselers$sex
```

```
## [1] "male"    "male"    "female"  "male"    "male"  
## [6] "male"    "male"    "female"  
## [ reached getOption("max.print") -- omitted 9992 entries]
```

```
# select sex variable using %>% select  
baselers %>% select(sex) %>% pull()
```

```
## [1] "male"    "male"    "female"  "male"    "male"  
## [6] "male"    "male"    "female"  
## [ reached getOption("max.print") -- omitted 9992 entries]
```

```
# Possible, but less pretty...  
baselers[['sex']]  
baselers[[2]]
```

Essentials: 10 basic R lessons

1. Everything is an object
2. <- creates/changes objects
3. Everything happens through functions
4. Functions have (default) arguments
5. Find help with ?
6. Data lives in data frames
7. **3 data types + factors**
8. formula and data specify a model
9. Use RStudio and projects
10. Use editor, shortcuts, auto-complete

```
# original sex vector  
baselers$sex
```

```
## [1] "male"  "male"  "female" "male"  "male"  
## [6] "male"  
## [ reached getOption("max.print") -- omitted 9994 entries]
```

```
# original sex vector  
as.factor(baselers$sex)
```

```
## [1] male  male  female male  male  male  
## [ reached getOption("max.print") -- omitted 9994 entries]  
## Levels: female male
```

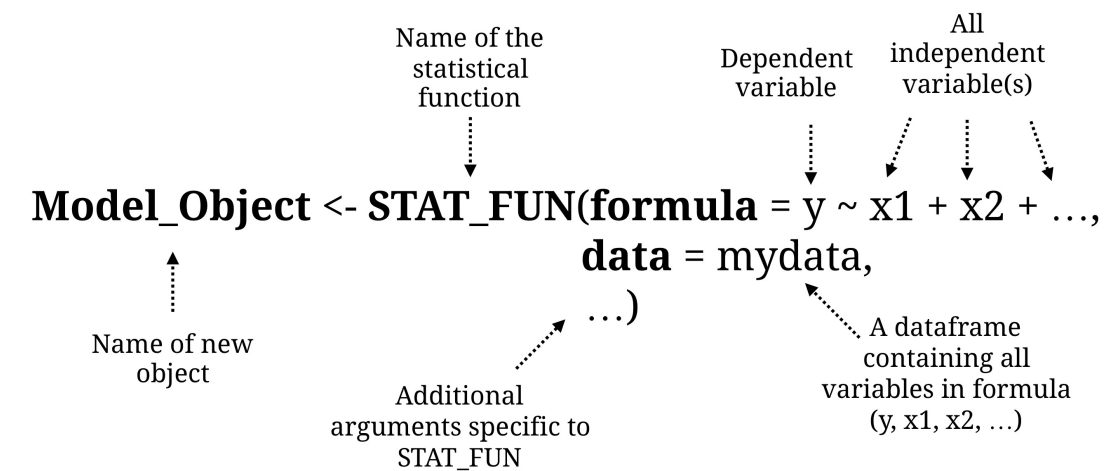
```
# original sex vector  
as.factor(baselers$weight)
```

```
## [1] 113.4 75.2 55.5 93.8 <NA> 67.4  
## [ reached getOption("max.print") -- omitted 9994 entries]  
## 719 Levels: 37.9 38.3 39.2 39.6 40.3 ... 125.4
```

Essentials: 10 basic R lessons

1. Everything is an object
2. `<-` creates/changes objects
3. Everything happens through functions
4. Functions have (default) arguments
5. Find help with `?`
6. Data lives in data frames
7. 3 data types + factors
8. **formula and data specify a model**
9. Use RStudio and projects
10. Use editor, shortcuts, auto-complete

```
# Run a regression and store result in my_lm
my_lm <- lm(formula = income ~ age + height,
            data = baselers)
```



Essentials: 10 basic R lessons

1. Everything is an object
2. `<-` creates/changes objects
3. Everything happens through functions
4. Functions have (default) arguments
5. Find help with `?`
6. Data lives in data frames
7. 3 data types + factors
8. **formula and data specify a model**
9. Use RStudio and projects
10. Use editor, shortcuts, auto-complete

Add variables using `+`

```
# Include multiple terms with +
my_lm <- lm(formula = income ~ age + height,
            data = baselers)
```

Include all variables using `formula = y ~ .`

```
# Use y ~ . to include ALL variables
my_lm <- lm(formula = income ~ .,
            data = baselers)
```

Subtract variables using `-`


```
# Use y ~ . to include ALL variables
my_lm <- lm(formula = income ~ . - id,
            data = baselers)
```

Essentials: 10 basic R lessons

1. Everything is an object
2. <- creates/changes objects
3. Everything happens through functions
4. Functions have (default) arguments
5. Find help with ?
6. Data lives in data frames
7. 4 types of data
8. formula and data specify a model
9. **Use RStudio and projects**
10. Use editor, shortcuts, auto-complete

Projects help...

save workspace and history • set project specific options •
access files • version control • etc.




Essentials: 10 basic R lessons

1. Everything is an object
2. `<-` creates/changes objects
3. Everything happens through functions
4. Functions have (default) arguments
5. Find help with ?
6. Data lives in data frames
7. 4 types of data
8. formula and data specify a model
9. **Use RStudio and projects**
10. Use editor, shortcuts, auto-complete

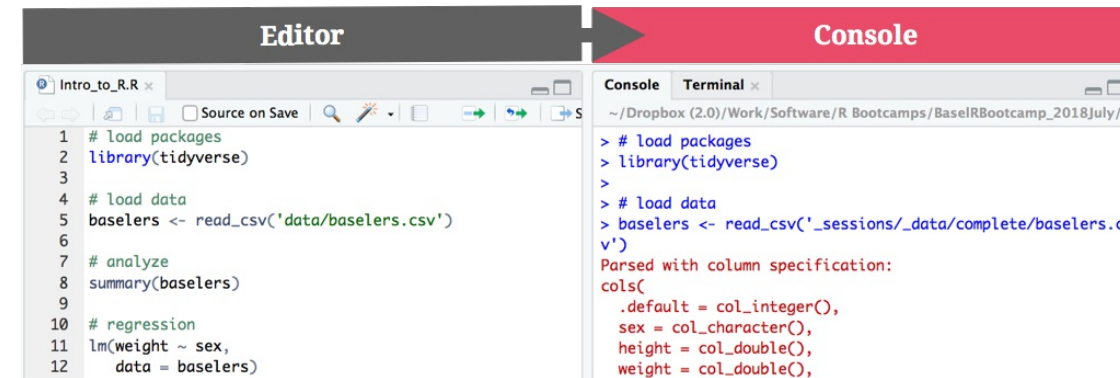
Folder structure

Complement projects by a **folder structure** appropriate for your project.



Essentials: 10 basic R lessons

1. Everything is an object
2. `<-` creates/changes objects
3. Everything happens through functions
4. Functions have (default) arguments
5. Find help with `?`
6. Data lives in data frames
7. 4 types of data
8. formula and data specify a model
9. Use RStudio and projects
10. **Use editor, shortcuts, auto-complete**



```
1 # load packages
2 library(tidyverse)
3
4 # load data
5 baselers <- read_csv('data/baselers.csv')
6
7 # analyze
8 summary(baselers)
9
10 # regression
11 lm(weight ~ sex,
12     data = baselers)
```

```
> # load packages
> library(tidyverse)
>
> # load data
> baselers <- read_csv('_sessions/_data/complete/baselers.csv')
Parsed with column specification:
cols(
  .default = col_integer(),
  sex = col_character(),
  height = col_double(),
  weight = col_double(),
```

Shortcut to **send to console**:

⌘/ctrl + ↵

Shortcut to **rerun chunk**:

⌘/ctrl + ⇧ + p

Essentials: 10 basic R lessons

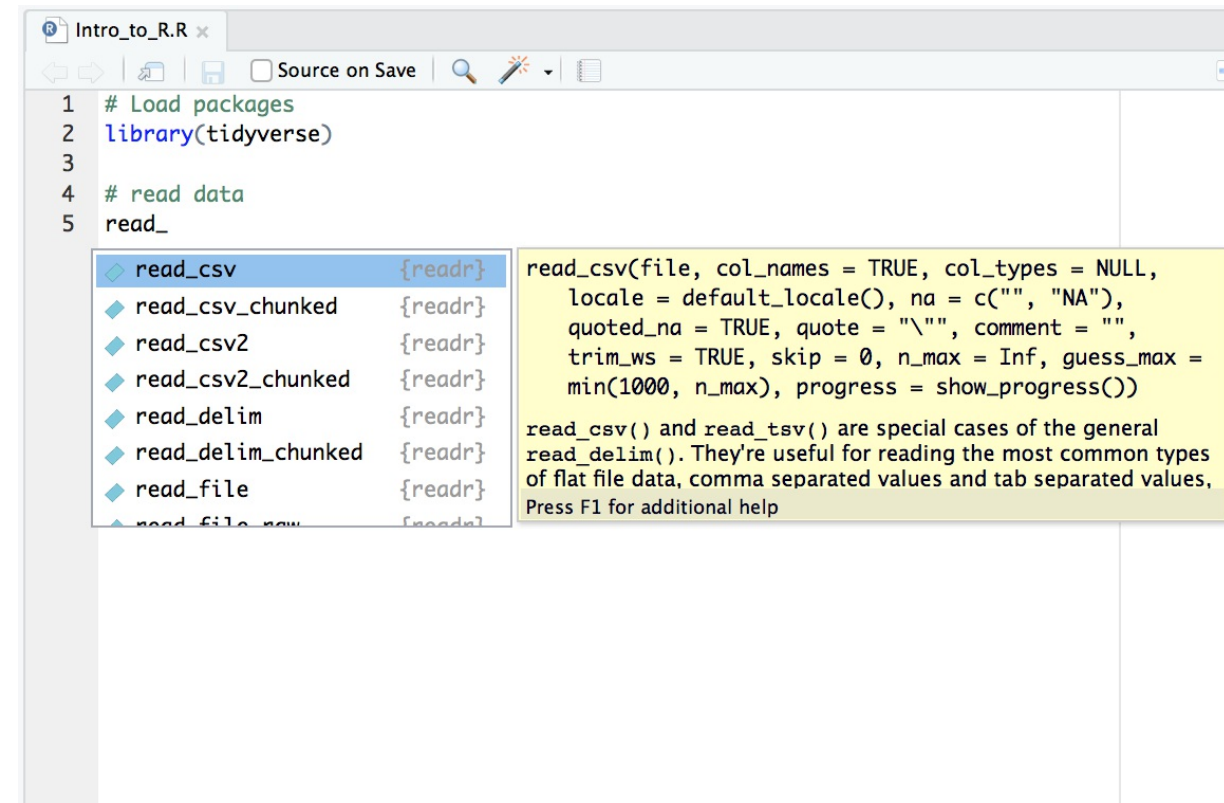
1. Everything is an object
2. <- creates/changes objects
3. Everything happens through functions
4. Functions have (default) arguments
5. Find help with ?
6. Data lives in data frames
7. 4 types of data
8. formula and data specify a model
9. Use RStudio and projects
10. **Use editor, shortcuts, auto-complete**

```
# import packages
library(tidyverse)
library(yarr)
library(lme4)

# import data
baselers <- read_delim("baselers.txt",
                      delim = '\t')
```

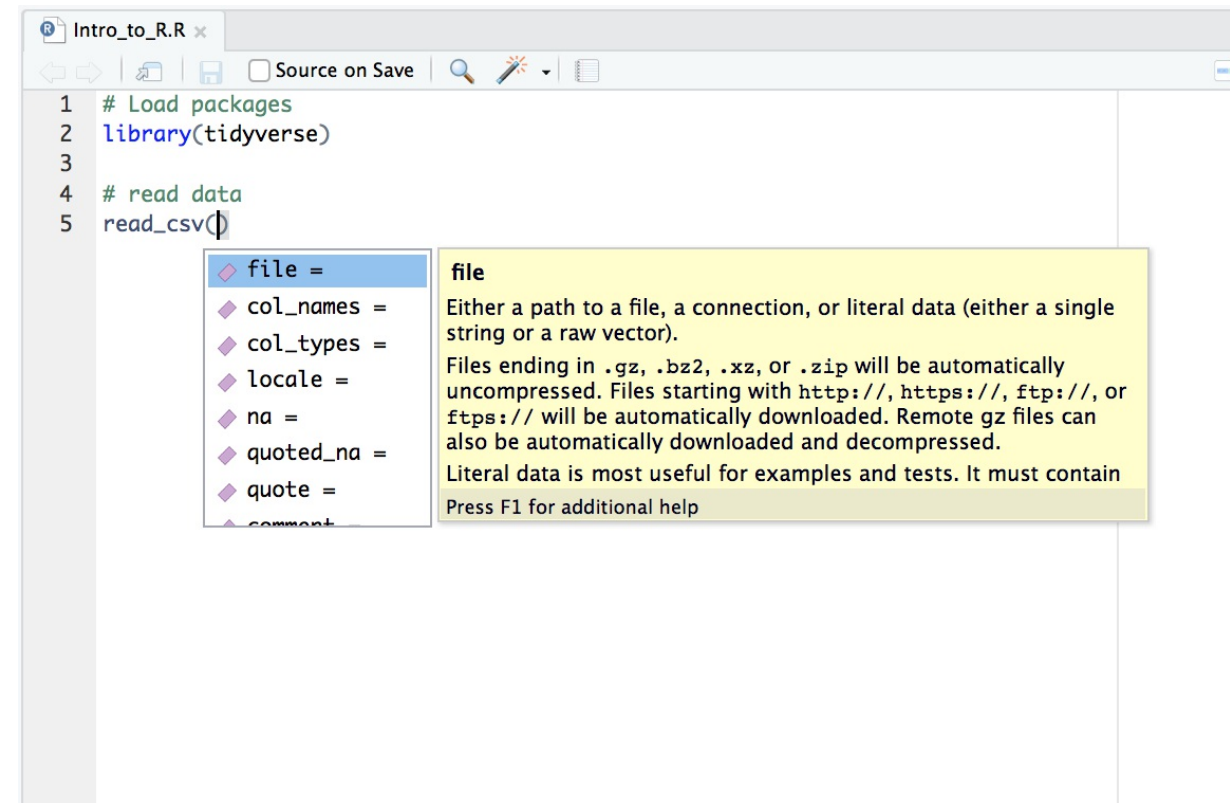
Essentials: 10 basic R lessons

1. Everything is an object
2. <- creates/changes objects
3. Everything happens through functions
4. Functions have (default) arguments
5. Find help with ?
6. Data lives in data frames
7. 4 types of data
8. formula and data specify a model
9. Use RStudio and projects
10. **Use editor, shortcuts, auto-complete**




Essentials: 10 basic R lessons

1. Everything is an object
2. <- creates/changes objects
3. Everything happens through functions
4. Functions have (default) arguments
5. Find help with ?
6. Data lives in data frames
7. 4 types of data
8. formula and data specify a model
9. Use RStudio and projects
10. **Use editor, shortcuts, auto-complete**



Essentials: 10 basic R lessons

1. Everything is an object
2. `<-` creates/changes objects
3. Everything happens through functions
4. Functions have (default) arguments
5. Find help with `?`
6. Data lives in data frames
7. 4 types of data
8. formula and data specify a model
9. Use RStudio and projects
10. **Use editor, shortcuts, auto-complete**



Download data sets

Interactive